

Python for Data Analytics

A series of white, concentric, rounded rectangular lines in the top right corner, creating a sense of depth and movement.

Module 2: Data structures and descriptive statistics





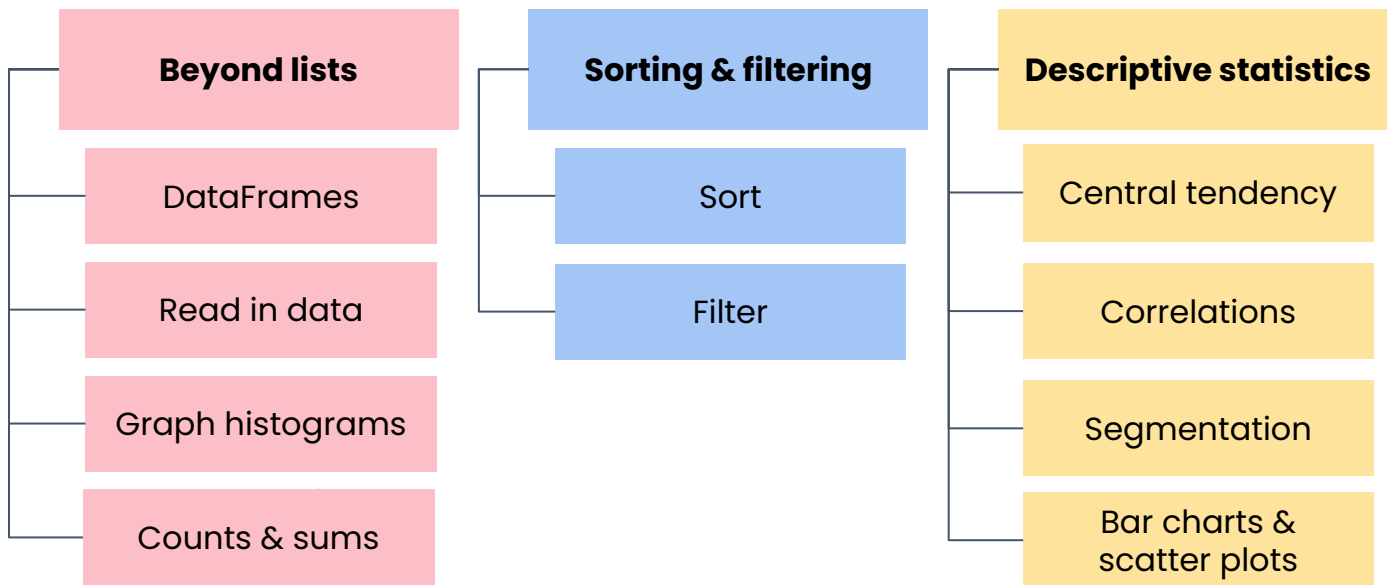
Data structures and descriptive statistics

Module 2 introduction

Module 2 outline



Coding habits





Data structures and descriptive statistics

Beyond lists

Data structures

- How data is arranged and organized

Lists

- | | |
|-----------------------|------------------------|
| ✓ Quick analysis | ✗ One dimensional |
| ✓ Collection of items | ✗ Too flexible |
| ✓ Action on items | ✗ Inefficient at scale |

Other data structures

- ✓ Two-dimensional
- ✓ Store multiple columns of data together

“What score did The Melrose Shrimp have?”

```
scores = [96, 91, 79, 93, 86]
```

✗ Missing key information

```
names = ["Beverly Falafel", "Pasta Roma", "The  
Melrose Shrimp", "Modern Eats", "Alferd's Coffee"]
```

```
scores = [96, 91, 79, 93, 86]
```

✗ This approach is clunky

✗ No inherent organization of rows and columns

Flexibility

✓ Can store multiple data types

✗ Doesn't have a built in way to add `int` and `string`

✗ Can hinder analysis if you introduce different types

- Each column should contain one type:
 - All ints
 - All strings
- Select a data structure that prevents introducing multiple types into a column

int string

→ `scores = [96, "91", 79, 93, "86"]` ✓

→ `print(sum(scores))` ✗

Error

→ `scores = [96, "orca", 79, 93, "narwhal"]` ✓

→ `print(sum(scores))` ✗

Error

Efficiency

Lists



- Introduce inefficiencies for very large datasets (i.e. data at scale)
- **for** loops are not the most efficient option at scale

Specialized data structures

- Make this type of operation much more efficient
- **Examples:** Dataframes, Series
- Perform the same operation on the entire data structure at once

```
A = 90  
for score in scores:  
    if score >= A:  
        print("A")
```

5 seconds



Vectorization: Performing an operation on the entire data structure at once



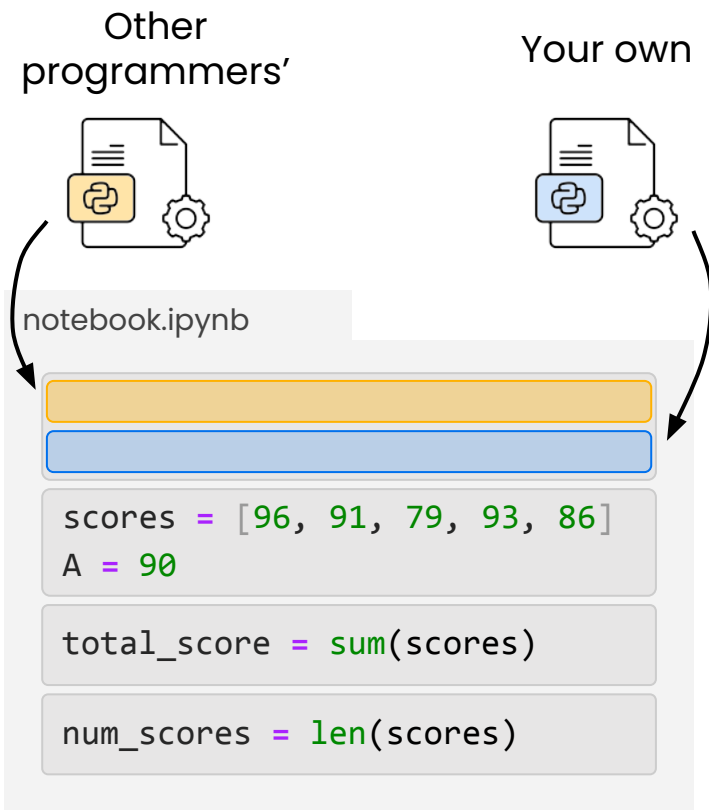
Data structures and descriptive statistics

Importing modules

Importing code

- Automatically have access to many built-in Python functions
- Borrow other code using:

```
import
```



Recap: Importing modules

A **module** is a file containing Python code that's already been written

1. Specific functions:

```
from helper_functions import get_restaurant_list
get_restaurant_list()
```

Labels: **module** (above `helper_functions`), **function_name** (above `get_restaurant_list`)

2. All functions:

```
import helper_functions
helper_functions.get_restaurant_list()
```

```
import helper_functions as hf
hf.get_restaurant_list()
```



helper_functions.py



Data structures and descriptive statistics

Pandas

pandas



- Extremely popular data science module
- Provides powerful data structures and functionality

DataFrame

Fiscal Year	Library	Operating expenses
1996	Avon	516583
1996	Plainville	49690
2005	Eastford	471864

Series

0	81003
1	72072
2	82698
3	96435

- Efficiently manipulate, clean, and analyze large datasets

- Provides function to:
 - Easily handle missing data
 - Sort, filter, and pivot your data
 - Merge and join datasets
 - Group and aggregate data
 - Perform time-series analysis
- Supports importing data from various file formats like:
 - CSV
 - Excel
 - Databases





Data structures and descriptive statistics

Reading CSV files
into Python

Scenario



You
Data Analyst



Task: Investigate characteristics of programmers



Goal: Report about new programmers:

- Demographics
- Descriptive statistics
- Visualizations



Data: Online survey of new programmers

1. Load into Python
2. Explore it
3. Characterize different features

Recap: Reading CSV files into a dataframe

Goal: Get CSV data loaded into your notebook to analyze it with Python

1. Know the name of the file

 Coder_analysis.ipynb

 survey_data.csv

2. Use `pd.read_csv` and with one argument: name of file as a string
3. Assign that data into a variable

Creates a **DataFrame** with the csv data

```
import pandas as pd
```

```
df = pd.read_csv("survey_data.csv")
```



```
data = pd.read_csv("survey_data.csv")
```

```
survey_data = pd.read_csv("survey_data.csv")
```

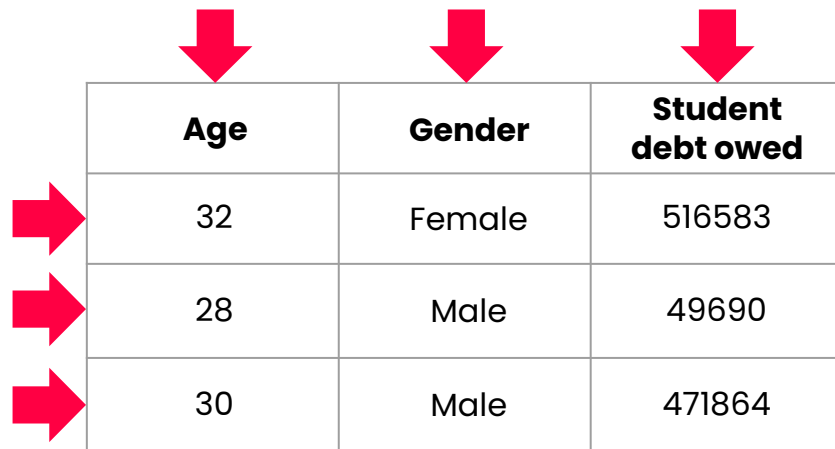


Data structures and descriptive statistics

DataFrames

DataFrame

- Represent a table or spreadsheet of data
- A type of data structure that represents data in **rows** and **columns**
- Similar to lists stored together, each representing a column



Age	Gender	Student debt owed
32	Female	516583
28	Male	49690
30	Male	471864

```
age = [32, 28, 30, ...]  
gender = ["Female", "Male", "Male", ...]  
student_debt_owed = [516583, 49690, 471864, ...]
```

Object type

- Store data that is more complex than numbers or true/false values.
- Pandas uses object type for texts to handle strings of any length
- When you see object:
 - Assume column contains text

Numbers:



- same amount of space

Text:



- "Great nail clippers!"



- Several paragraphs



Data structures and descriptive statistics

Attributes and methods

Attributes and methods

Attributes:

`df.columns`

- Do not need parentheses
- Something DataFrame **has**
- Computer needs to fetch it
- No calculations needed

`df.dtypes`

Method:

`df.info()`



- Something DataFrame can **do**
- **Generate** a summary of its data
- Fancy word for function

Commands:

`sum()`

`max()`

`len()`



- Have parentheses because these are **actions**
- Computer adds up all values

Analogy

Yourself

Computer

Attributes

- “What’s your name?”
- “What’s your age?”

- Commonly accessed, fundamental info

Method

- “How many days has it been since you visited a park?”

- More complex questions
- Generating summary data
- Most operations

Big picture: You’ll need to differentiate between a DataFrame’s attributes and methods.

- When you’re accessing an **attribute**, like columns or dtypes, you don’t need to use parentheses
- When you’re using a **method**, you will use parentheses, and often arguments to the method as well



Data structures and descriptive statistics

Selecting columns

Selection

- Choosing specific part of data:
 - Rows
 - Columns
 - Individual values

- Selection with **lists**:

```
scores[0]
```

Selects the first item in scores list

Selecting columns

1 Select a single column with a line of code:

Returns a Pandas Series, a 1 dimensional data structure

```
languages = df["LanguagesAtHome"]
```



2 Select multiple columns by:

1. Saving the columns in a list
2. Placing that list inside the brackets

```
columns = ["CountryOfResidence", "LanguageAtHome"]  
country_columns = df[columns]
```




Data structures and descriptive statistics

Counts, sums, &
histograms

Recap: Counts, sum, & histograms

1 Use the `.count()` method directly on:

- A dataframe to see the number of non-null values for each column:

```
df.count()
```

- A series:

```
df["Age"].count()
```

2 Use `.hist()` to plot histogram showing the distribution of a numerical feature

```
df["Income"].hist()
```

3 Use the `.sum()` method to add up all the values in a numeric column

```
df["HoursSpentLearningToCode"].sum()
```



Data structures and descriptive statistics

Sorting

Recap: Sorting

- Sort DataFrames using:

```
df.sort_values(by = "Age", ascending = True)
```

Column to
sort by

- Save data frame in new variable

```
df_sorted_by_age = df.sort_values(by = "Age", ascending = True)
```

Named
arguments



Data structures and descriptive statistics

Sorting by multiple
columns

Recap: Sorting by multiple columns

List 1: Columns to sort by

List 2: Values for ascending



In both cases, sort in descending order

Save DataFrame into a new variable,
sorted_by_age_and_hours

1
↓
columns = ["Age", "HoursSpentLearningToCode"]

2
↓
order = [False, False]

sorted_by_age_and_hours = df.sort_values(by=columns, ascending=order)

Taking DataFrame and sorting by values:

- Sort by columns in the columns list
- Specifying to sort ascending or descending



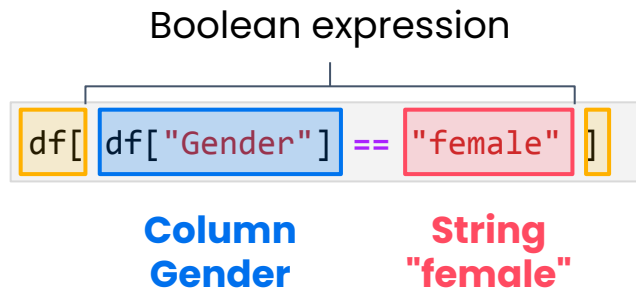
Data structures and descriptive statistics

Filtering

Recap: Filtering

- To select rows that match condition:
 - Select from data frame with a boolean expression
- The boolean expression inside the brackets can be:
 - Equal to ==
 - Greater than >
 - Less than <

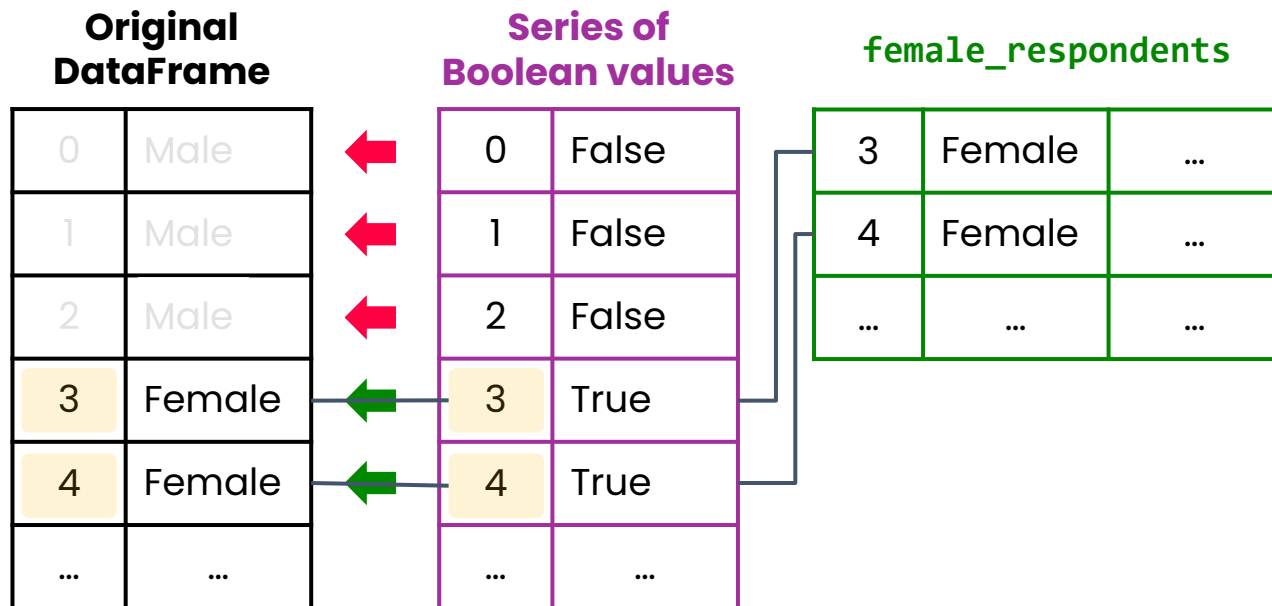
Example: Select only responses where the survey taker answered "female" for gender



Recap: Filtering

```
female_respondents = df[df["Gender"] == "female"]
```

- **True** if row is female
- **False** in any other case





Data structures and descriptive statistics

Filtering by multiple
conditions

Recap: Filtering by multiple conditions

Filter by multiple conditions using:

- **And operator:** select rows that meet both conditions

```
female_above_30 = df[(df['gender'] == 'female') & (df['age'] >= 30)]
```



- **Or operator:** select rows that meet at least one conditions

```
female_or_above_30 = df[(df['gender'] == 'female') | (df['age'] >= 30)]
```





Data structures and descriptive statistics

Selecting rows

Indices

- **Numerical** indices

Each row has a unique identifier that's a number, starting with 0

- Pandas gives option to assign custom indices like:




Student IDs



Country codes


- Assign strings or other types

DataFrame



	Age	Hours
0	19	8
1	21	23
2	18	4
3	26	11
4	18	12

Series



0	Male
1	Male
2	Male
3	Female
4	Female

Recap: Selecting rows

- Access rows based on indices:

```
indexed_df = df.set_index("ID")  
indexed_df.loc[25447]
```

- Select row based on its position:

First row is at index 0

```
indexed_df.iloc[0]
```

- To select a slice of rows:

```
indexed_df.iloc[1000:1006]
```

- Select all the rows from a, including a, to b-1, not b


```
indexed_df.iloc[a:b]
```

a	1000	✓
	1001	✓
	1002	✓
	1003	✓
	1004	✓
b - 1	1005	✓
b	1006	✗

A note on slicing

```
df.iloc[1000:1006]
```

- Counterintuitive that you type in the value when slicing, but don't get that value
- **Main benefit:** Length of the slice is always end minus start



$$1006 - 1000 = 6$$



Length of this slice is 6.



If slice included last value, the length would be 7



Data structures and descriptive statistics

Central tendency,
variability, and skewness

Recap: Central tendency, variability, and skewness

- Use `.describe()` on a series or DataFrame to calculate descriptive statistics about numerical features:

```
hours.describe()
```

→ count	14942.000000
→ mean	15.323317
→ median	10.000000
std	14.274867
min	0.000000
25%	5.000000
50%	10.000000
75%	20.000000
max	100.000000

- Calculate these values individually using the different pandas methods:

`.mean()` `.median()`

`.skew()`

`.max()` `.min()`

`.std()` `.var()`

} No arguments

→ `.quantile()`

```
hours.quantile(0.5)
```

```
hours.quantile([0.25, 0.5, 0.75])
```



Data structures and descriptive statistics

Categorical data

Recap: Categorical data

With categorical data you can:

- Use the `.value_counts()` method, which returns:
 - A Series in descending order of the number of occurrences
- Create a column chart by:
 - Saving the result of `value_counts`
 - Using `.plot()` method with the named argument `kind = "bar"`

```
df["Gender"].value_counts()
```

Gender

male	10766
female	2840
genderqueer	66
agender	38
trans	36

Indices

Values

```
counts_of_gender = df["Gender"].value_counts()  
counts_of_gender.plot(kind="bar")
```



Data structures and descriptive statistics

Correlation

Recap: Correlation

To create a scatter plot:

- Use `.plot()` method on a data frame
- "kind" named argument is "scatter"
- "x" named argument is feature on x axis
- "y" named argument is feature on y axis

```
df.plot(kind="scatter", x="Income", y="MoneySpentLearningToCode")
```

To calculate correlations:

- Select only the numerical columns
- Use the `.corr()` method on the subset

```
selected_columns = df[columns]  
selected_columns.corr()
```

	Age	NumberOfChildren	MoneySpentLearningToCode	MonthsSpentProgramming	Income
Age	1.000000	0.240286	0.098985	0.223237	0.259090
NumberOfChildren	0.240286	1.000000	-0.009486	0.048951	0.099751
MoneySpentLearningToCode	0.098985	-0.009486	1.000000	0.086157	0.078810
MonthsSpentProgramming	0.223237	0.048951	0.086157	1.000000	0.286328
Income	0.259090	0.099751	0.078810	0.286328	1.000000



Data structures and descriptive statistics

Segmentation by
one feature

Recap: Segmentation by one feature

Groupby **doesn't calculate** until you take the next step to perform some calculations

For segmentation, use the `.groupby()` method:

```
df.groupby("NumberofChildren")
```

pandas.core.groupby.generic.DataFrameGroupBy 

Select columns from the groupby:

```
df.groupby("NumberofChildren")["HoursSpentLearningToCode"]
```

Perform computations:

```
df.groupby("NumberofChildren")["HoursSpentLearningToCode"].count()
```

... sum, or mean, or other descriptive statistics

NumberOfChildren

1.0	1050
2.0	938
3.0	342
4.0	109
5.0	37
6.0	10
7.0	1
8.0	1
9.0	2
10.0	2
12.0	1
15.0	0
18.0	1



Data structures and descriptive statistics

Segmentation by
multiple features

Recap: Segmentation by multiple features

Use to group data by more than one feature:

```
df.pivot_table(index = "NumberOfChildren",  
               columns = "IsSoftwareDev",  
               values = "HoursSpentLearningToCode" )  
               aggfunc = "sum")
```

- Use the aggfunc named argument to specify a different summarization


Some options: sum, count, std, max, and min

- Provide a list for aggfunc argument, which gives a pivot table containing summarization functions

Features to group by

Outcome of interest

Default: values will be summarized using the mean



```
...aggfunc = ["min", "median", "max"])
```